

FIT3080 – Artificial Intelligence

Assignment Project Exam Help

Adversarial Search

<https://powcoder.com>

Add WeChat powcoder

Algorithms

Announcements

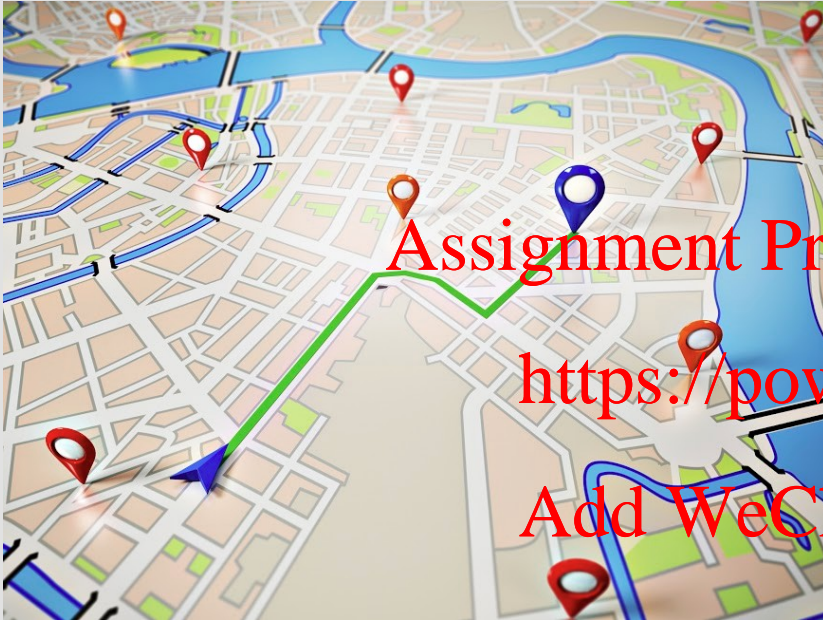
- Assignment 1 is out!
 - Modelling and solving problems using search
 - 4 parts (mix of conceptual and practical tasks)
 - Check Moodle for the specification
 - Due date: September 7, 11:55pm (Melbourne)



Last week

- Informed tree/graph search
 - Expand 'most promising' nodes first,
 - Using evaluation function: $f(n) = g(n) + h(n)$
- Heuristics <https://powcoder.com>
 - What are they
 - Properties & guarantees
 - Effectiveness
- Bounded suboptimal search

Last week




Route planning




Puzzles

Assignment Project Exam Help
<https://powcoder.com>
Add WeChat powcoder

Last week



Route



Coordination problems



zles



This week: games!



Tic-Tac-Toe

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



This week: games!



Tic-Tac-Toe



Chess

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



This week: games!



Tic-Tac-Toe



Chess

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Games with an adversary

This week: games!



Tic-Tac-Toe



Chess

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Games with an adversary

Man vs. Machine!



This week: games!



Chinook vs. Tinsley (1993/4)

Watch Jonathan Schaffer wonderful account:

<https://www.youtube.com/watch?v=VWqtNS9pmOI>



This week: games!



Assignment Project Exam Help

<https://powcoder.com>

Deep Blue vs. Gary Kasparov (1996/7)

Add WeChat powcoder

A short recap from the BBC:

<https://www.youtube.com/watch?v=KF6sLCeBj0s>



MONASH University
Information Technology

This week: games!



Deep Blue vs. Gary Kasparov (1996/7)

A short recap from the BBC:
<https://www.youtube.com/watch?v=KF6sLCeBj0s>



This week: games!



AlphaGo vs. Lee Sedol (2016)

There's a wonderful documentary about this match:

<https://www.youtube.com/watch?v=WXuK6gekU1Y>



Types of games

- Perfect information (deterministic, full obsv.)
 - Chess, Go, noughts-and-crosses (tic-tac-toe), draughts (checkers), etc.
- Imperfect information (stochastic, partially obsv.)
 - Poker, Texas hold'em (variant of poker), bridge, backgammon, etc.
- n-Players (e.g. $n = 2$ for chess)
- Sequential (vs. simultaneous)
- Zero-sum (vs non-zero sum games)

Our focus: 2p, sequential, perfect information, zero sum



Game Trees

- Start state: initial configuration of the board
- Players are MAX and MIN
 - A position favourable to MAX has a value > 0 (winning often ∞)
 - A position favourable to MIN has a value ≤ 0 (winning often $-\infty$)
- Each player chooses the most promising move for them
- Each move adds new nodes to the frontier
- The set of nodes generated for one player is called a ply
- Each node has an associated utility for each player
- Leaves of the tree are called terminal nodes

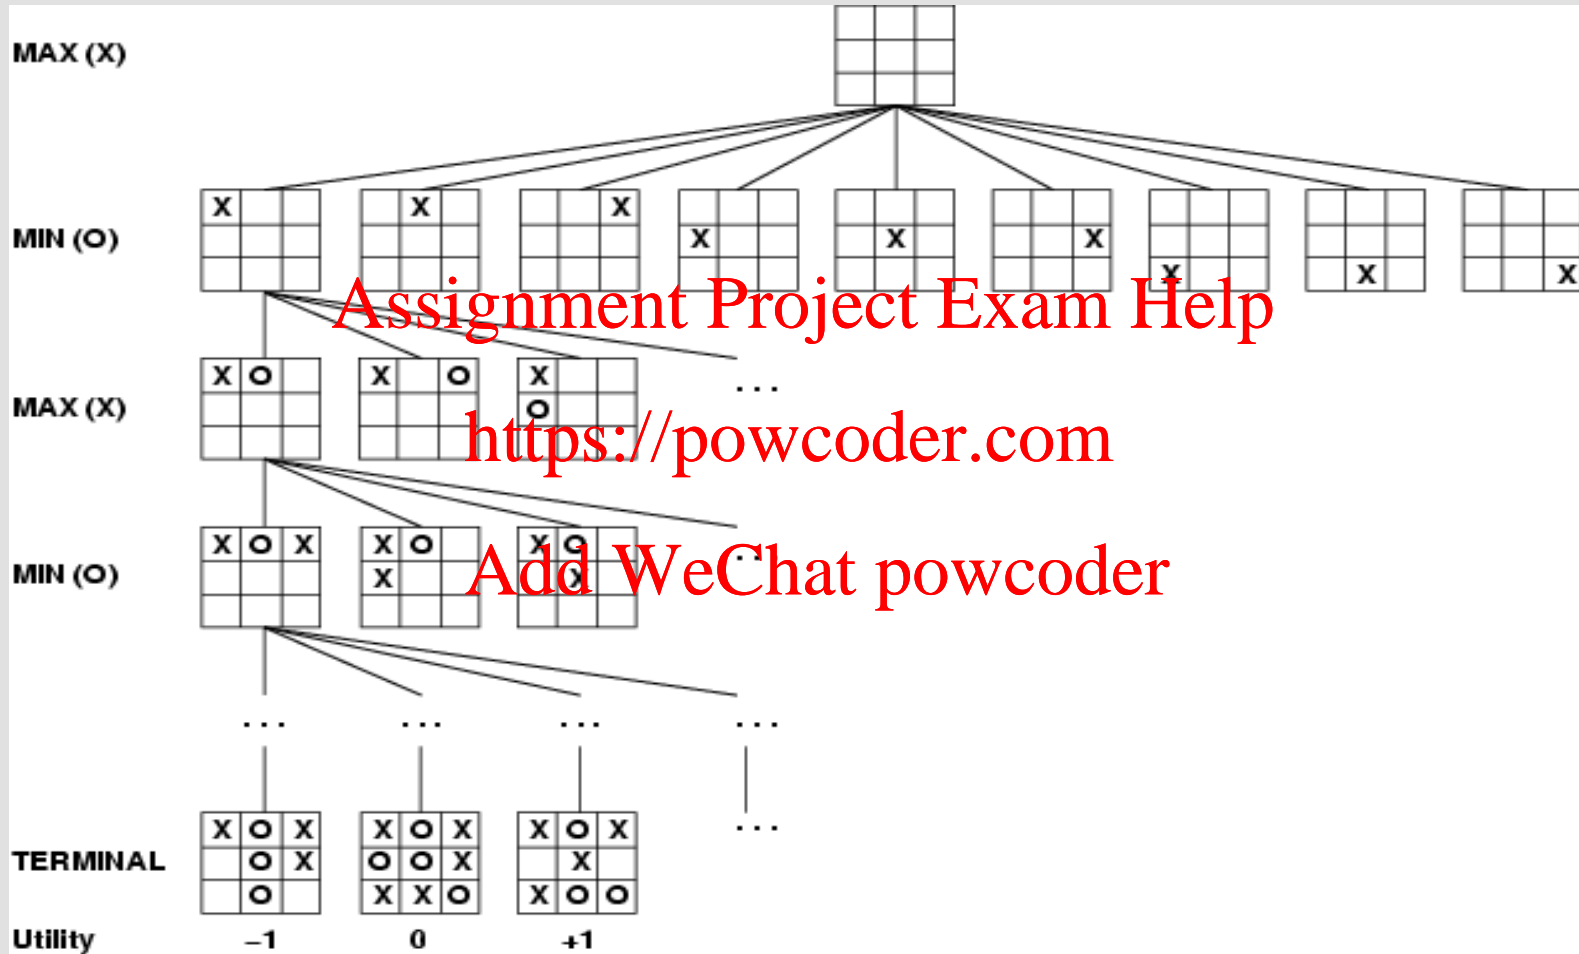
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Game Tree (2-player, Deterministic, Turns)



Solving Game Trees

A solution is a game-tree where we can label the root node with an outcome (win, loss, draw, or a utility value)

- Each move available to the opponent is a subproblem (another game)
- All subproblems need to be solved before the parent node can be considered solved

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

A strategy for MAX (resp. MIN) is a solution-tree which contains exactly one successor for each MAX (resp. MIN) node.

A winning strategy for MAX (resp. MIN) is a solution-tree where every terminal is a WIN (resp. LOSS, resp. max utility value)

Solving Game Trees

Goal: find a winning strategy for MAX

- For all nodes representing a game situation where it is MIN's move next, show that MAX can win from **every** position to which MIN might move
- For all nodes representing a game situation where it is MAX's move next, show that MAX can win from **just one** position to which MAX might move

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder




Games versus Search Problems

- **Unpredictable opponent:**
 - **Must specify a move for every possible opponent reply**
Assignment Project Exam Help
- **Time limits:** <https://powcoder.com>
 - **not all games can be searched to the end**
Add WeChat powcoder
 - **find a good first move**

Search ideas

- If MAX were to choose among tip nodes (or available nodes), she would take the node with the largest value
- If MIN were to choose among tip nodes (or available nodes), she would take the node with the smallest value
- Choose a move to the position with the highest minimax value: best achievable payoff against best play
- E.g., 2-ply game:



Minimax Algorithm

```
function MINIMAX-DECISION(state) returns an action  
  return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$ 
```

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

Assignment Project Exam Help

<https://powcoder.com>

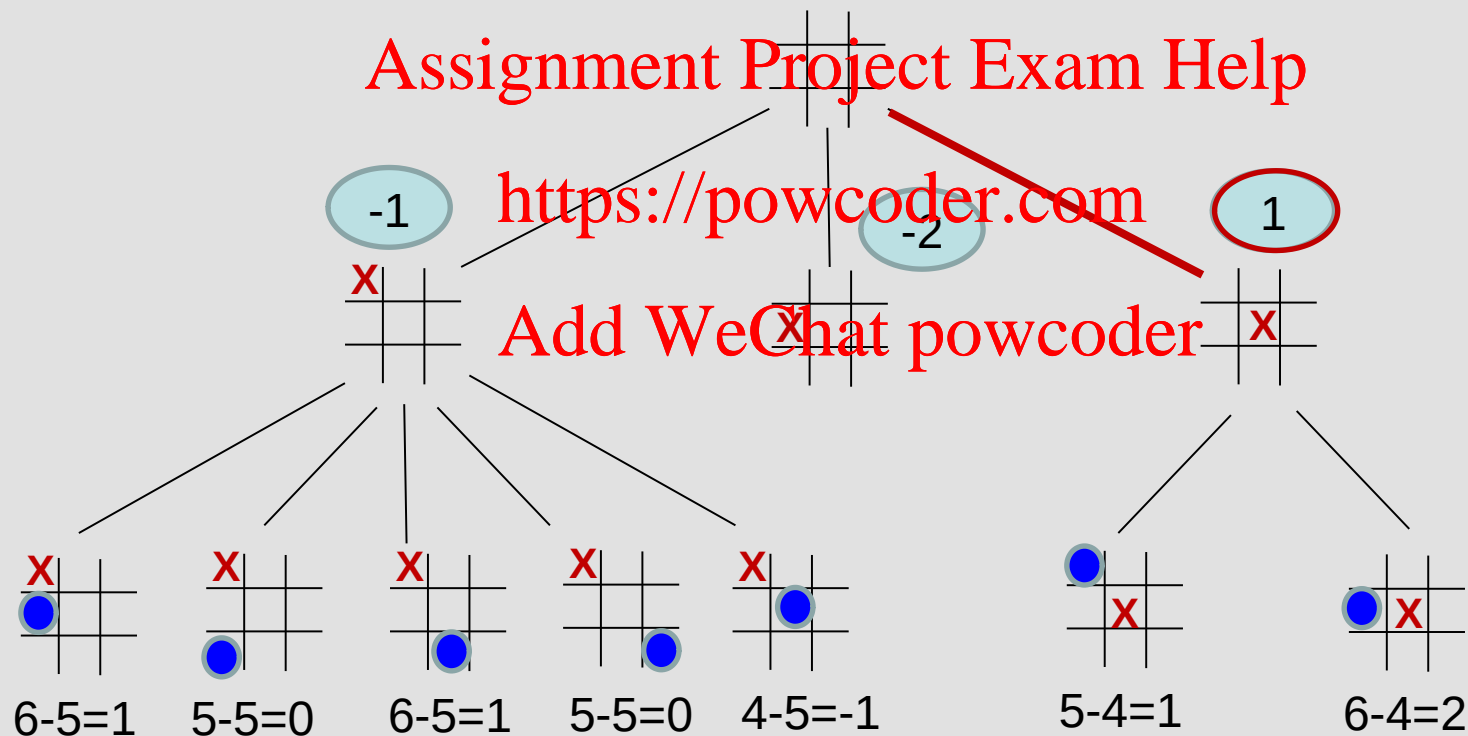
Add WeChat powcoder



Minimax Example: Tic-Tac-Toe

Evaluation function:

{ # rows, columns, diagonals available to MAX –
of rows, columns, diagonals available to MIN }



Properties of Minimax

Complete Depth First Exploration:

all paths to depth m

Complete? Yes (if tree is finite)

Optimal? Yes (against an optimal opponent)

Time complexity? $O(b^m)$

Space complexity? $O(bm)$ (depth first exploration)

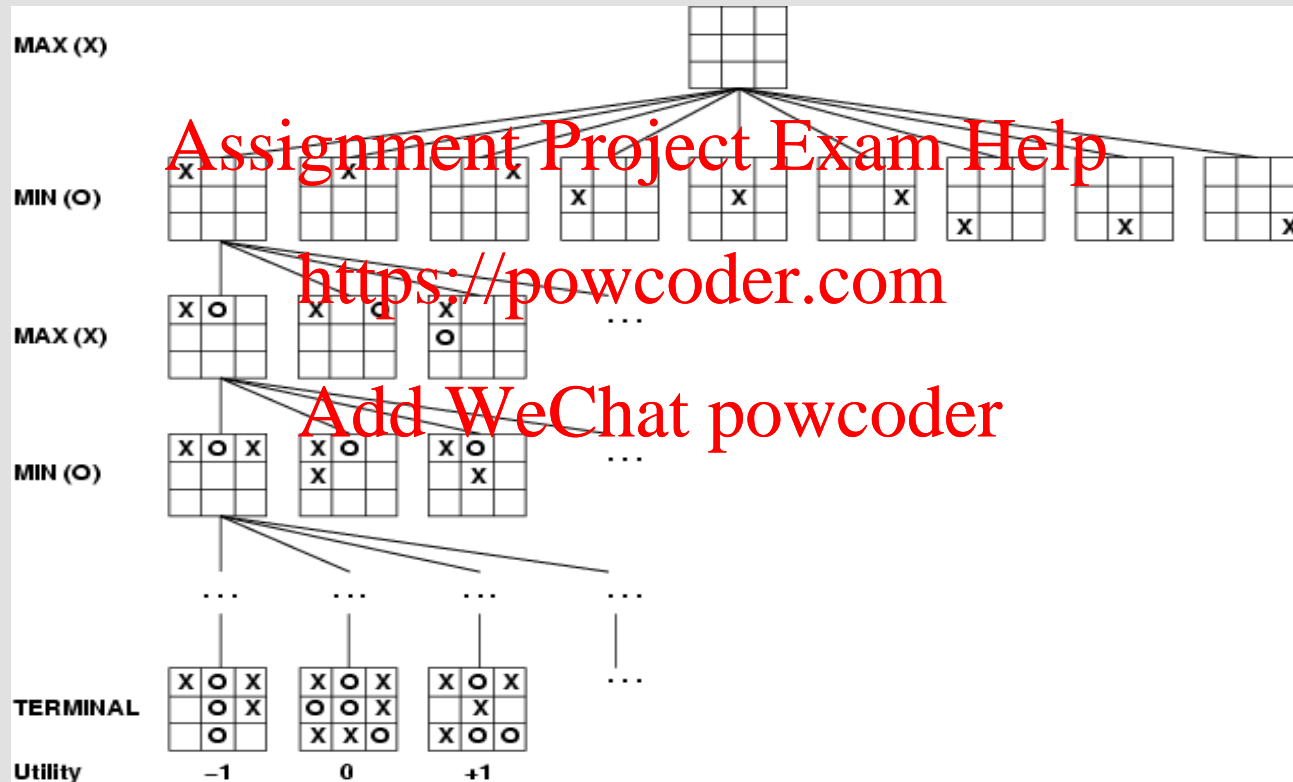
For chess, $b \approx 35$, $m \approx 100$ for “reasonable” games

→ exact solution completely infeasible



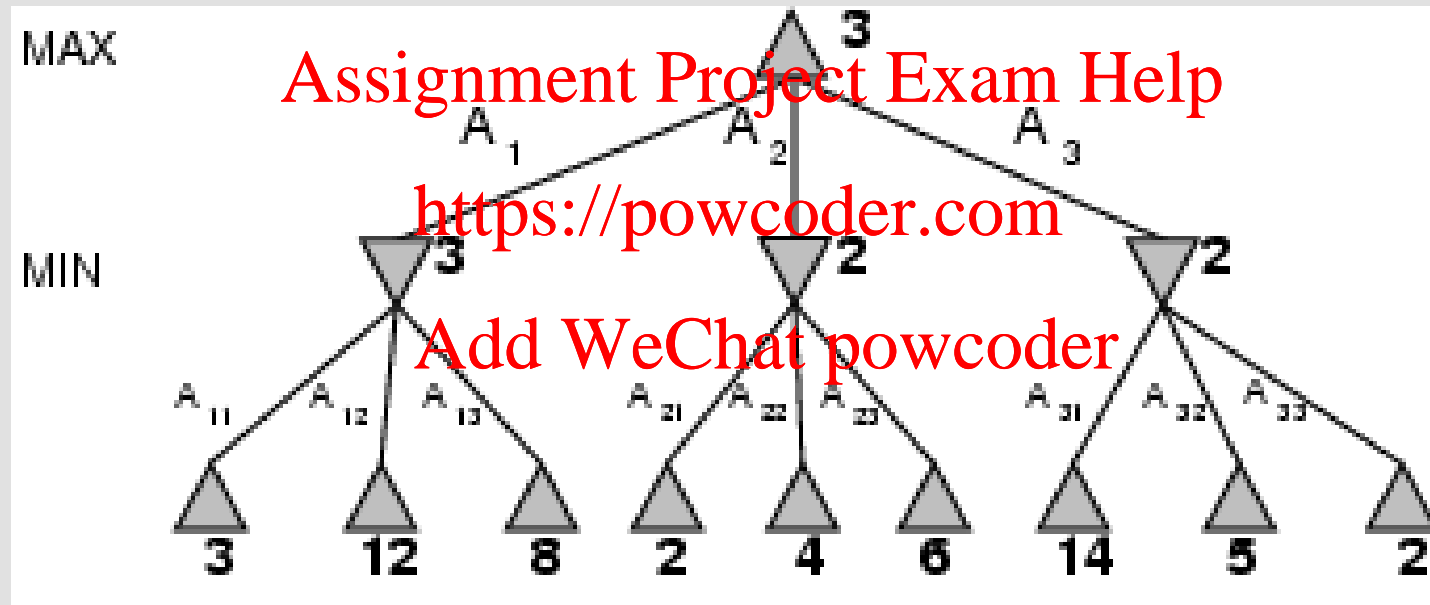
Pruning the Search Space

We can improve performance by eliminating symmetric moves



Pruning the Search Space

We can improve performance by eliminating redundant nodes



Definitions: α and β Values

- **α -value of a MAX node – current largest final backed-up value of its successors**
 - α -value is the lower bound for a MAX backed-up value
 - **Assignment Project Exam Help**
- **β -value of a MIN node – current smallest final backed-up value of its successors**
 - β -value is the upper bound for a MIN backed-up value

<https://powcoder.com>

Add WeChat powcoder



α - β Procedure

- **Rules for discontinuing the search:**
 - **α cut-off:** search can be discontinued below any **MIN** node having a β -value \leq **α -value** of **any** of its MAX node ancestors
 - The final backed-up value of this MIN node is set to its β -value
 - **β cut-off:** search can be discontinued below any **MAX** node having an α -value \geq **β -value** of **any** of its MIN node ancestors
 - The final backed-up value of this MAX node is set to its α -value

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Termination Condition

- All the successors of the start node are given final backed-up values
- The best first move is that which creates the successor with the highest backed-up value

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Minimax Algorithm (again)

```
function MINIMAX-DECISION(state) returns an action  
  return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$ 
```

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



The α - β Algorithm (I)

```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   $\alpha$  cut-off  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   $\beta$  cut-off  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```

α = the value of the best (i.e., **highest-value**) choice we have found so far at any choice point along the path for MAX.

β = the value of the best (i.e., **lowest-value**) choice we have found so far at any choice point along the path for MIN.

Assignment Project Exam Help

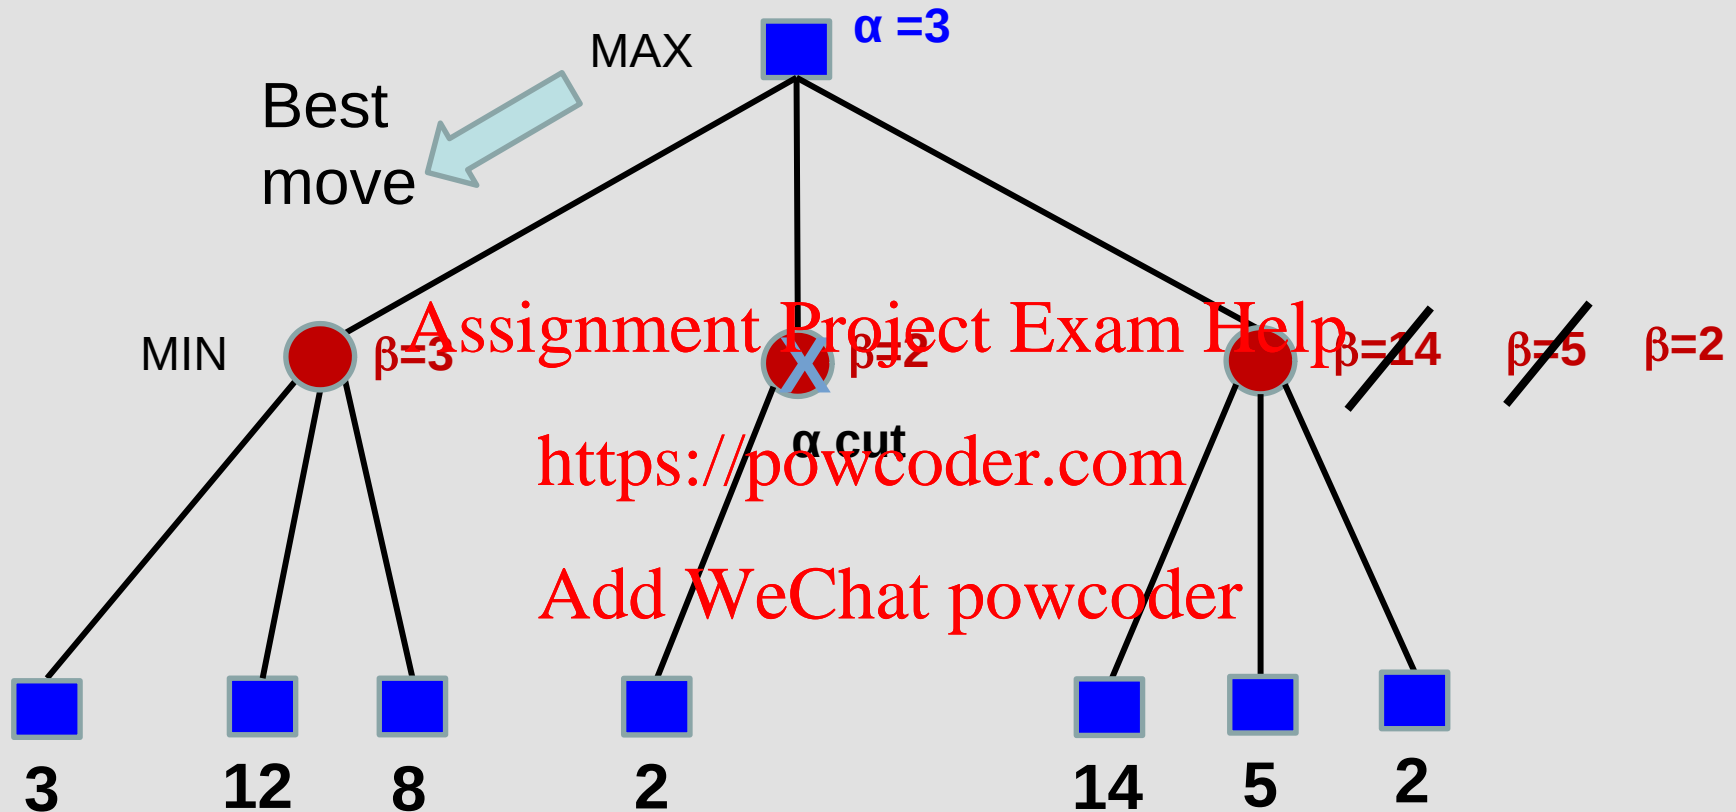
<https://powcoder.com>

Add WeChat powcoder

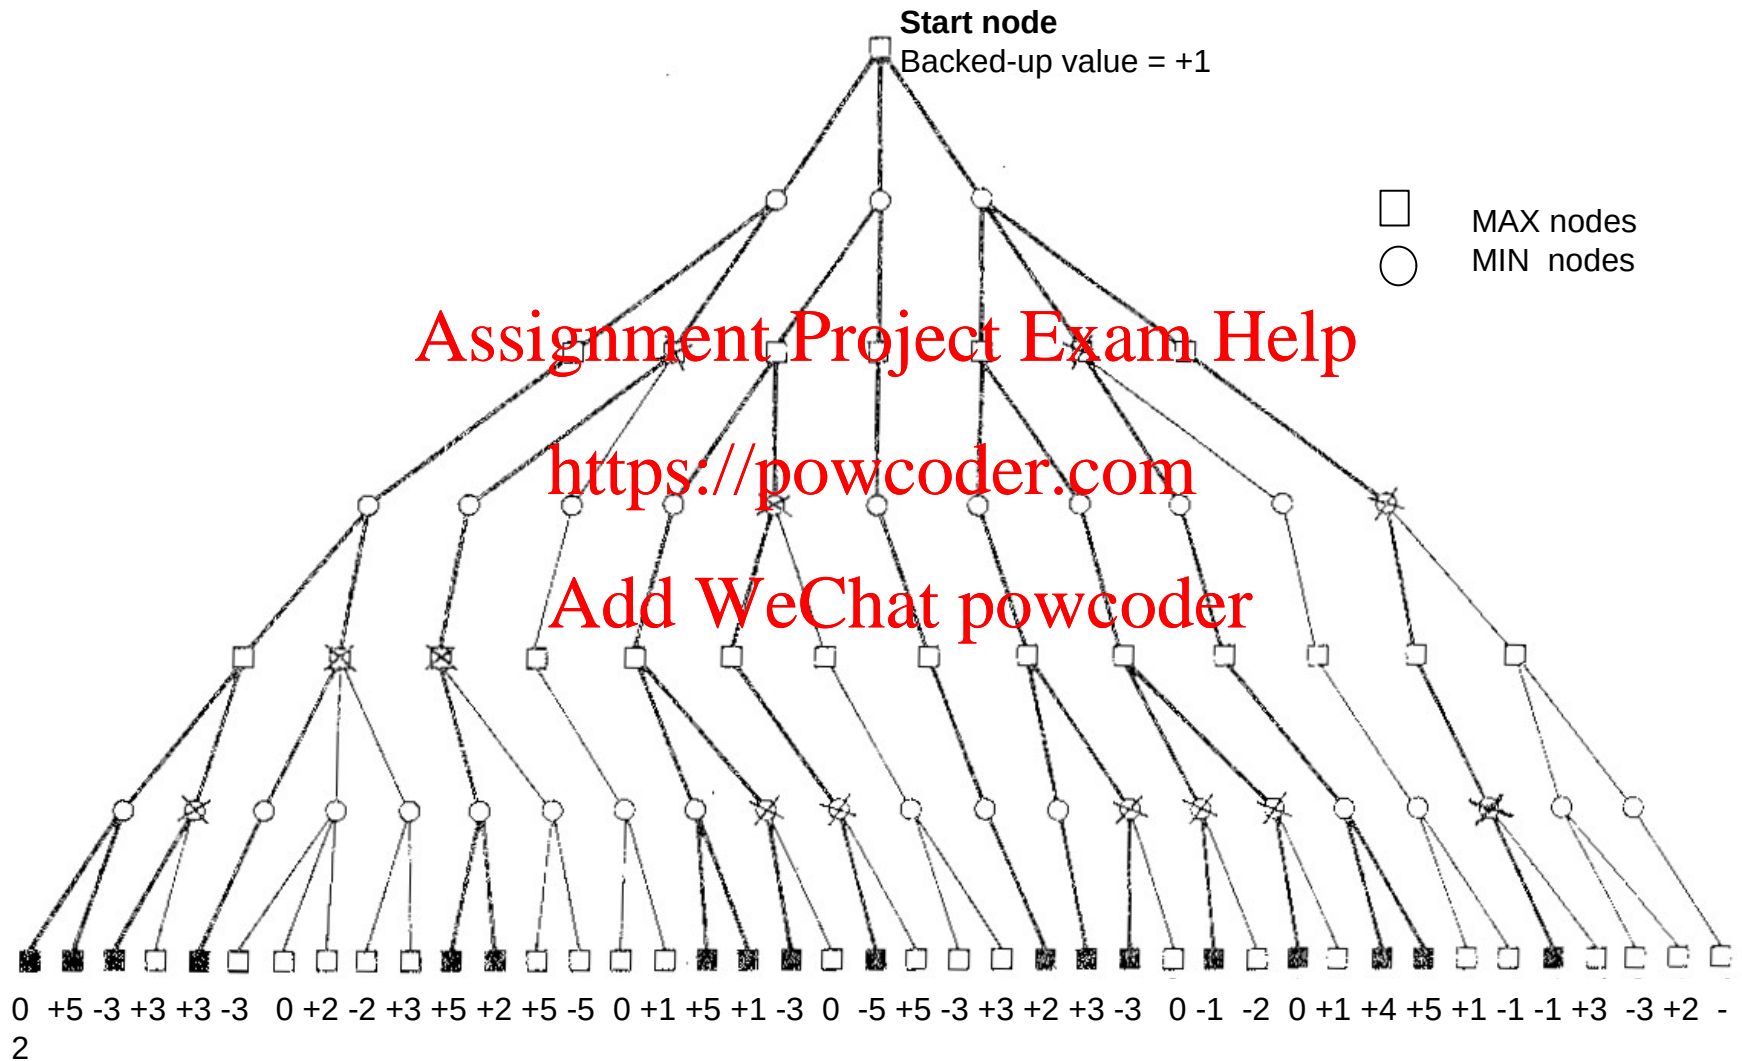
α cut-off




α - β Pruning – Example



α - β Pruning – Large Example



α - β Pruning – Part of Large Example



Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder



α - β Pruning – Part of Large Example

Does the order of the moves matter?



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Move Ordering

The effectiveness of the $\alpha\beta$ algorithm depends on the order in which states are examined

- With perfect ordering:
- time complexity = $O(\sqrt{b}^n) = O(b^{n/2})$
- i.e., depth of search can be doubled

Add WeChat powcoder

Adding dynamic ordering schemes brings us close to the theoretical limit



Heuristics for Games

Suppose we have 100 secs per move, and we explore 10^4 nodes/sec $\rightarrow 10^6$ nodes per move

- Still might not reach a terminal node!

Assignment Project Exam Help

Idea: treat the frontier nodes as terminal states

<https://powcoder.com>
Add WeChat powcoder

$$\text{H-MINIMAX}(s, d) = \begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN}. \end{cases}$$

Evaluation Functions

Some ideas for generating these.

From experience:

- **Material values** in chess (1 = pawn, 3 = bishop, ... 9 = Queen)
- **State features:**
- $f_1(s) = (\# \text{ of white queens}) - (\# \text{ of black queens})$
- **Combinations of features** (e.g., linear weighted sum)

$$\text{Eval}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s),$$

From preprocessing:

- End-game databases



How do we search?

Depth limited search (fixed cost)

Depth-First Iterative Deepening (anytime)

Assignment Project Exam Help

Beam search (forward pruning)

<https://powcoder.com>


Don't search at all (use table lookups)

Add WeChat powcoder



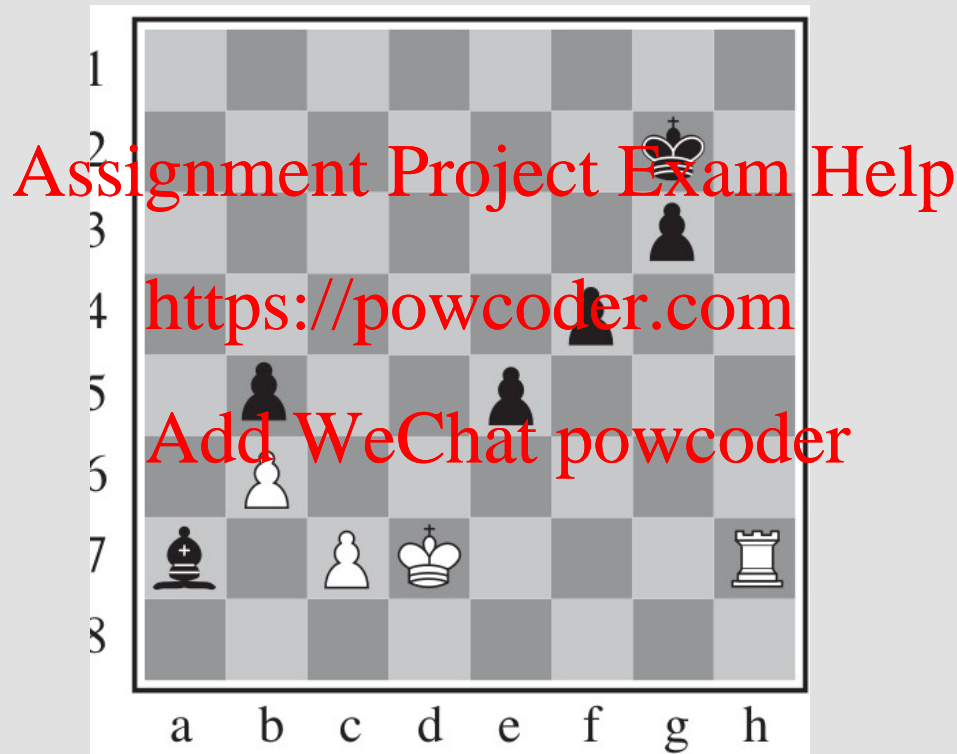
Search considerations

Potentially misleading evaluations



Search considerations

Potentially misleading evaluations



Heuristic AB Tree Search

Starting from a fixed board position: (state, player)

- 1) Run AB up to some limit (DFS, or DFID)
- 2) Treat the frontier nodes s' as terminal states
- 3) Compute a score using $\text{EVAL}(s', p)$
- 4) Back up the values
- 5) Pick the best move
- 6) Repeat from the new board position

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Types of Search

- Type 1:
wide, but shallow
- Type 2:
narrow, but deep
- **Historically practitioners preferred Type 1 Searches.**
- **More recently, Type 2 approaches (and Type1/2 combinations) have emerged.**

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Monte Carlo Simulation

Use sampling to obtain estimates of utility at a given state

- Starting from a given position (s, p)
- “Roll out” the game until a terminal state
- Fixed policy determines the moves at each ply (fast!)
- Repeat many times to obtain an “average” utility

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Monte Carlo Tree Search

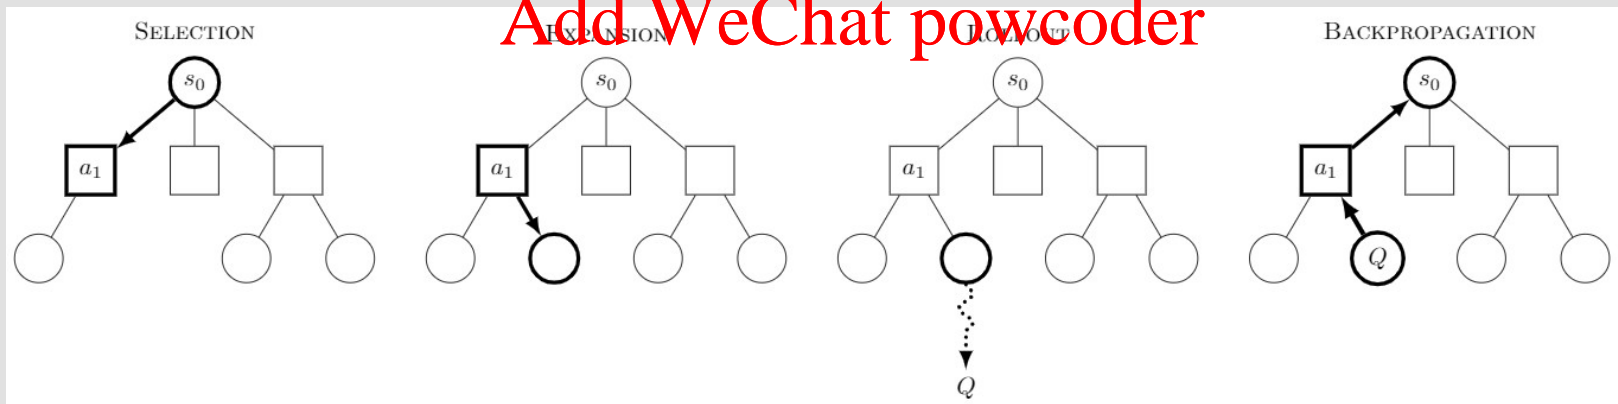
Four main ingredients:

- 1) Selection: pick a node from the frontier
- 2) Expand: generate successors to extend the frontier
- 3) Simulate games from the newly added node
- 4) Propagate the results

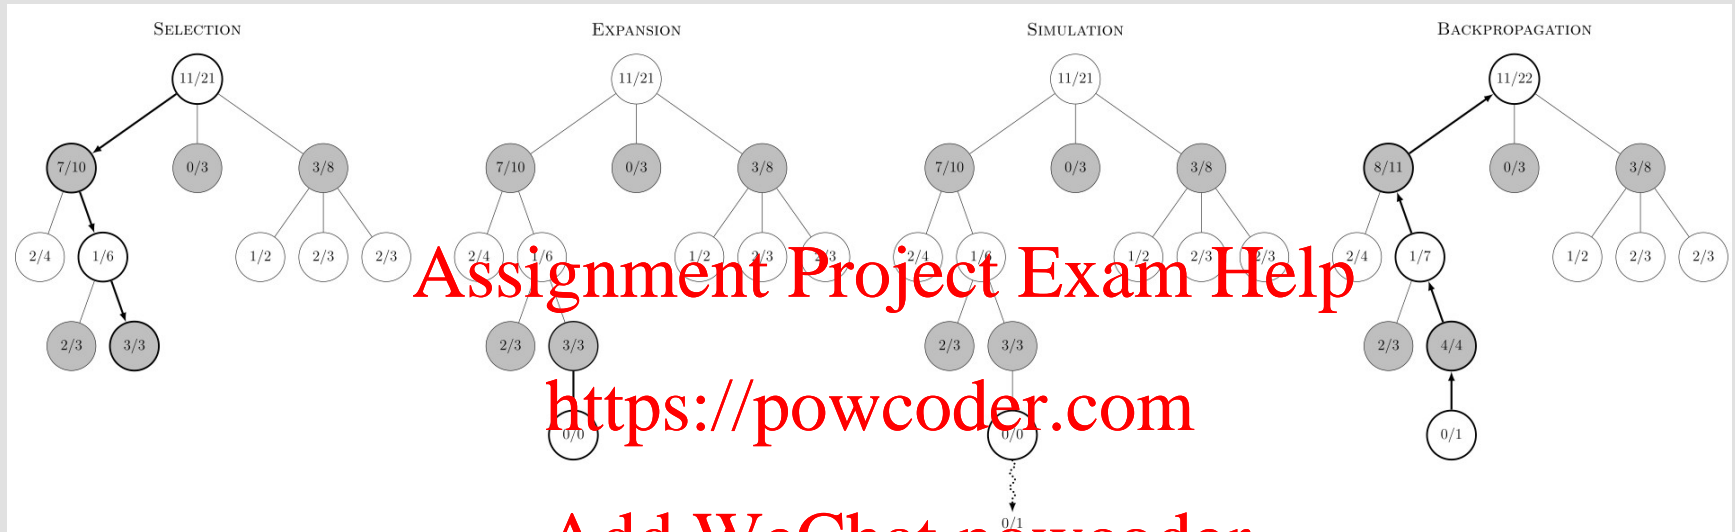
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Monte Carlo Tree Search



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Instead of always maximising average utility, we can bias the choices: **exploration vs exploitation**



Deterministic Games in Practice

- **Checkers:** Chinook defeated the world champion in an abbreviated game in 1990. It uses $\alpha\beta$ search combined with a pre-computed database defining perfect play for 39 trillion endgame positions.
- **Chess:** Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 30 billion positions per move (200 million per second), normally searching to depth 14, and extending the search up to depth 40 for promising options. Heuristics reduce the EBF to about 3.
- **Othello:** In 1997, a computer defeated the world champion 6-0. Humans are no match for computers.
- **Go:** $b \sim 19 \times 19 = 361$, which is too large for $\alpha\beta$. In 2016, AlphaGo, which uses Deep Learning, beat the world champion 4-1.
- AlphaGo also used Monte Carlo Tree Search (MTCS)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Summary: Adversarial Search

Games illustrate important points about AI

- **Perfection is unattainable → must approximate**
- **It causes one to think about which problems are worth thinking about**

Assignment Project Exam Help

<https://powcoder.com>

The basic ideas discussed today generalise:

Add WeChat powcoder

- **>2 players → more ply and more costs**
- **Stochastic or hidden information**
→ **probabilistic reasoning (expectimax)**



Next week – Logical Agents

- **Propositional logic**

- **Inference**

- **Resolution**

- **Satisfiability**

Assignment Project Exam Help

- **First Order Logic**

<https://powcoder.com>

- **Logic and agents**

Add WeChat powcoder

